Advanced Introduction to Machine Learning CMU-10715

Principal Component Analysis

Barnabás Póczos

ML MACHINE LEARNING DEPARTMENT

Carnegie Mellon.
School of Computer Science

Contents

❑Motivation

❑PCA algorithms

❑Applications

Some of these slides are taken from

- Karl Booksh Research group
- Tom Mitchell
- Ron Parr

Motivation

PCA Applications

- Data Visualization
- Data Compression
- Noise Reduction

Data Visualization

Example:

- Given 53 blood and urine samples (features) from 65 people.

- How can we visualize the measurements?

Data Visualization

- Matrix format (65x53)

	H-WBC	H-RBC	H-Hgb	H-Hct	H-MCV	H-MCH	H-MCHC
A1	8.0000	4.8200	14.1000	41.0000	85.0000	29.0000	34.0000
A2	7.3000	5.0200	14.7000	43.0000	86.0000	29.0000	34.0000
A3	4.3000	4.4800	14.1000	41.0000	91.0000	32.0000	35.0000
A4	7.5000	4.4700	14.9000	45.0000	101.0000	33.0000	33.0000
A5	7.3000	5.5200	15.4000	46.0000	84.0000	28.0000	33.0000
A6	6.9000	4.8600	16.0000	47.0000	97.0000	33.0000	34.0000
A7	7.8000	4.6800	14.7000	43.0000	92.0000	31.0000	34.0000
A8	8.6000	4.8200	15.8000	42.0000	88.0000	33.0000	37.0000
A9	5.1000	4.7100	14.0000	43.0000	92.0000	30.0000	32.0000

Instances

Features

Difficult to see the correlations between the features…

Data Visualization

- Spectral format (65 curves, one for each person)

Difficult to compare the different patients...

Data Visualization

- Spectral format (53 pictures, one for each feature)

Difficult to see the correlations between the features…

Data Visualization

Bi-variate

Tri-variate

How can we visualize the other variables???

... difficult to see in 4 or higher dimensional spaces...

Data Visualization

- Is there a representation better than the coordinate axes?

- Is it really necessary to show all the 53 dimensions?
  - ... what if there are strong correlations between the features?

- How could we find
  the *smallest* subspace of the 53-D space that
  keeps the *most information* about the original data?

- A solution: **Principal Component Analysis**

PCA Algorithms

Principal Component Analysis

PCA:

Orthogonal projection of the data onto a lower-dimension linear space that...

❑ maximizes variance of projected data (purple line)

❑ minimizes the mean squared distance between
  - data point and
 - projections (sum of blue lines)

Principal Component Analysis

Idea:

❑ Given data points in a d-dimensional space, project them into a lower dimensional space while preserving as much information as possible.

- Find best planar approximation of 3D data
- Find best 12-D approximation of 10^4-D data

❑ In particular, choose projection that minimizes *squared error* in reconstructing the original data.

Principal Component Analysis

Properties:

❑ **PCA Vectors** originate from the center of mass.

❑ Principal component #1: points in the direction of the **largest variance**.

❑ Each subsequent principal component
- is **orthogonal** to the previous ones, and
- points in the directions of the **largest variance of the residual subspace**

2D Gaussian dataset

# 1st PCA axis

2nd PCA axis

PCA algorithm I (sequential)

Given the **centered** data $\{\mathbf{x}_1, ..., \mathbf{x}_m\}$, compute the principal vectors:

$$\mathbf{w}_1 = \arg\max_{\|\mathbf{w}\|=1} \frac{1}{m} \sum_{i=1}^{m} \{(\mathbf{w}^T \mathbf{x}_i)^2\} \qquad \text{1}^{st} \text{ PCA vector}$$

To find $\mathbf{w}_1$, maximize the variance of projection of $\mathbf{x}$

$$\mathbf{w}_2 = \arg\max_{\|\mathbf{w}\|=1} \frac{1}{m} \sum_{i=1}^{m} \{[\mathbf{w}^T (\underbrace{\mathbf{x}_i - \mathbf{w}_1 \mathbf{w}_1^T \mathbf{x}_i})]^2\} \qquad \text{2}^{nd} \text{ PCA vector}$$

$\mathbf{x'}$ PCA reconstruction

To find $\mathbf{w}_2$, we maximize the **variance** of the projection in the **residual** subspace

$\mathbf{x'}=\mathbf{w}_1(\mathbf{w}_1^\top\mathbf{x})$

Given $\mathbf{w_1}, \ldots, \mathbf{w_{k-1}}$, we calculate $\mathbf{w_k}$ principal vector as before:

Maximize the variance of projection of $\mathbf{x}$

$$\mathbf{w}_k = \arg\max_{\|\mathbf{w}\|=1} \frac{1}{m} \sum_{i=1}^{m} \{ [\mathbf{w}^T (\mathbf{x}_i - \underbrace{\sum_{j=1}^{k-1} \mathbf{w}_j \mathbf{w}_j^T \mathbf{x}_i})]^2 \}$$

k^{th} PCA vector

$\mathbf{x'}$ PCA reconstruction

We maximize the variance of the projection in the residual subspace

$\mathbf{w}$

$\mathbf{x}$

$\mathbf{w_1}(\mathbf{w}_1^T\mathbf{x})$ $\mathbf{w_1}$

$\mathbf{w_2}(\mathbf{w}_2^T\mathbf{x})$

$\mathbf{x'} = \mathbf{w}_1(\mathbf{w}_1^T\mathbf{x}) + \mathbf{w}_2(\mathbf{w}_2^T\mathbf{x})$

$\mathbf{w_2}$

PCA algorithm II (sample covariance matrix)

- Given data $\{\mathbf{x}_1, \ldots, \mathbf{x}_m\}$, compute covariance matrix Σ

$$\Sigma = \frac{1}{m} \sum_{i=1}^{m} (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \qquad \text{where} \qquad \bar{\mathbf{x}} = \frac{1}{m} \sum_{i=1}^{m} \mathbf{x}_i$$

- **PCA** basis vectors = the eigenvectors of Σ

- Larger eigenvalue $\Rightarrow$ more important eigenvectors

PCA algorithm($\mathbf{X}$, k): top k eigenvalues/eigenvectors

     % $\underline{\mathbf{X}}$ = N × m data matrix,
     % … each data point $\mathbf{x}_i$ = column vector, i=1..m

- $\underline{\mathbf{x}} = \dfrac{1}{m} \sum_{i=1}^{m} \mathbf{x}_i$

- $\mathbf{X} \leftarrow$ subtract mean $\underline{\mathbf{x}}$ from each column vector $\mathbf{x}_i$ in $\underline{\mathbf{X}}$

- $\Sigma \leftarrow \mathbf{X}\mathbf{X}^{\mathsf{T}}$ … covariance matrix of $\mathbf{X}$

- $\{ \lambda_i, \mathbf{u}_i \}_{i=1..N}$ = eigenvectors/eigenvalues of Σ
  ... $\lambda_1 \geq \lambda_2 \geq \ldots \geq \lambda_N$

- Return $\{ \lambda_i, \mathbf{u}_i \}_{i=1..k}$
 % top k PCA components

21

Finding Eigenvectors with Power Iteration

First eigenvector:

Power iteration 1: v=Sigma*v;

$\qquad$ v=v/sqrt(v'*v);

$\qquad$ $\Rightarrow$ v_{PCA1}

Second eigenvector:

Power iteration 2:

$\lambda = v_{PCA1}{}^T * Sigma * v_{PCA1}$

$Sigma_2 = Sigma - \lambda * v_{PCA1} * v_{PCA1}{}^T;$

$\qquad$ $v = Sigma_2 * v;$

$\qquad$ v=v/sqrt(v'*v);

$\qquad$ $\Rightarrow$ v_{PCA2}

Power Iteration: 1

Singular Value Decomposition of the **centered** data matrix **X**.

$$\mathbf{X} = [\mathbf{x}_1, \ldots, \mathbf{x}_m] \in \mathbb{R}^{N \times m},$$

m: number of instances,
N: dimension

$$\mathbf{X}_{\text{features} \times \text{samples}} = \mathbf{USV}^{\mathsf{T}}$$

$$\mathbf{X} = \mathbf{U} \quad \mathbf{S} \quad \mathbf{V}^{\mathsf{T}}$$

U: significant / noise

S: sig. / noise

V^T: significant / noise

samples

24

PCA algorithm III

- **Columns of U**
  - the principal vectors, $\{\ \mathbf{u}^{(1)}, ..., \mathbf{u}^{(k)}\ \}$
  - orthogonal and has unit norm – so $U^T U = I$
  - Can reconstruct the data using linear combinations of $\{\ \mathbf{u}^{(1)}, ..., \mathbf{u}^{(k)}\ \}$

- **Matrix S**
  - Diagonal
  - Shows importance of each eigenvector

- **Columns of V^T**
  - The coefficients for reconstructing the samples

Applications

Face Recognition

❑ Want to identify specific person, based on facial image

❑ Robust to glasses, lighting,…

⇒ Can't just use the given 256 x 256 pixels

Applying PCA: Eigenfaces

Method A: Build a PCA subspace for *each person* and check which subspace can reconstruct the test image the best

Method B: Build one PCA database for the *whole dataset* and then classify based on the weights.

Applying PCA: Eigenfaces

❑ Example data set: Images of faces
  - Eigenface approach
    [Turk & Pentland], [Sirovich & Kirby]

❑ Each face **x** is …

  - 256×256 values (luminance at location)
  - **x** in $\Re^{256 \times 256}$ (view as 64K dim vector)

❑ Form **X** = [**x**$_1$, …, **x**$_m$] **centered** data mtx

❑ Compute $\Sigma = \mathbf{XX}^\top$

❑ Problem: Σ is 64K $\times$ 64K … HUGE!!!

$\mathbf{x}_1, …, \mathbf{x}_m$

$\mathbf{X} =$

256 x 256 real values

m faces

Computational Complexity

❑ Suppose m instances, each of size N
  • Eigenfaces: m=500 faces, each of size N=64K
❑ Given N×N covariance matrix Σ, can compute
  • all N eigenvectors/eigenvalues in $O(N^3)$
  • first k eigenvectors/eigenvalues in $O(k N^2)$

❑ If N=64K, this is EXPENSIVE!

A Clever Workaround

- Note that $m \ll 64K$
- Use $\mathbf{L} = \mathbf{X^T X}$ instead of $\Sigma = \mathbf{X X^T}$
- If $\mathbf{v}$ is eigenvector of $\mathbf{L}$, then $\mathbf{Xv}$ is eigenvector of Σ

Proof:
$$\mathbf{L}\ \mathbf{v} = \gamma\ \mathbf{v}$$
$$\mathbf{X^T X}\ \mathbf{v} = \gamma\ \mathbf{v}$$
$$\mathbf{X\ (X^T X\ v)}\ =\ \mathbf{X}(\gamma\ \mathbf{v}) = \gamma\ \mathbf{Xv}$$
$$\mathbf{(XX^T)X\ v}\ =\ \gamma\ \mathbf{(Xv)}$$
$$\Sigma\ \mathbf{(Xv)}\ =\ \gamma\ \mathbf{(Xv)}$$

$\mathbf{x}_1, \ldots, \mathbf{x}_m$

$\mathbf{X} =$

256 x 256 real values

m faces

31

Principle Components (Method B)

Shortcomings

❑ Requires carefully controlled data:
- All faces centered in frame
- Same size
- Some sensitivity to angle

❑ Method is completely knowledge free
- (sometimes this is good!)
- Doesn't know that faces are wrapped around 3D objects (heads)
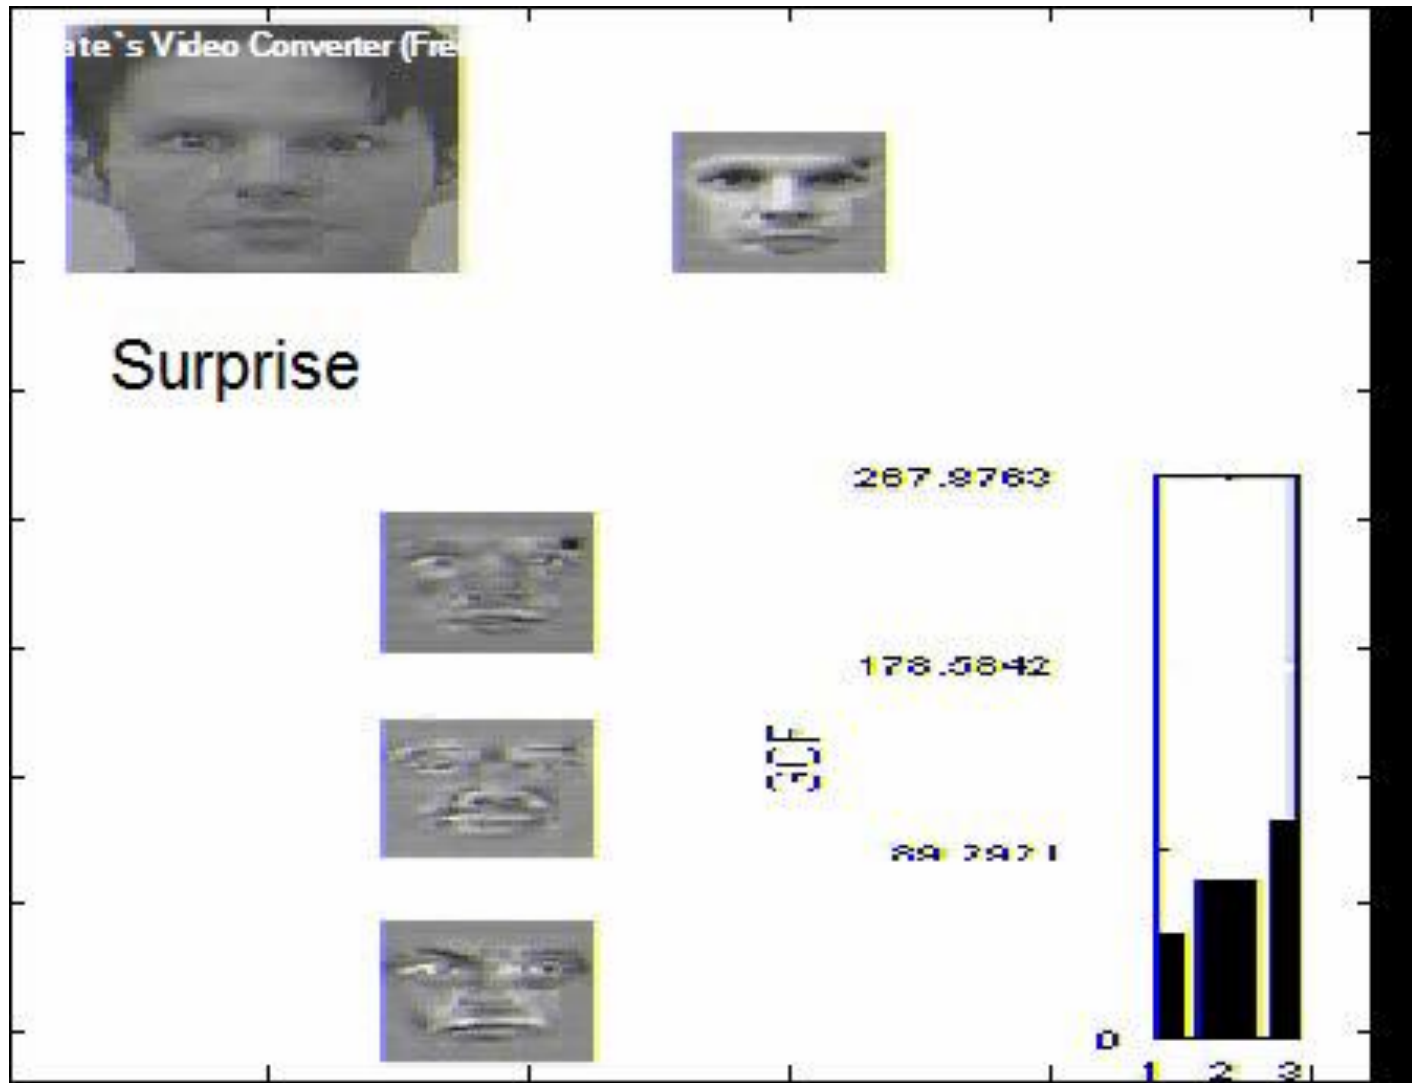- Makes no effort to preserve class distinctions

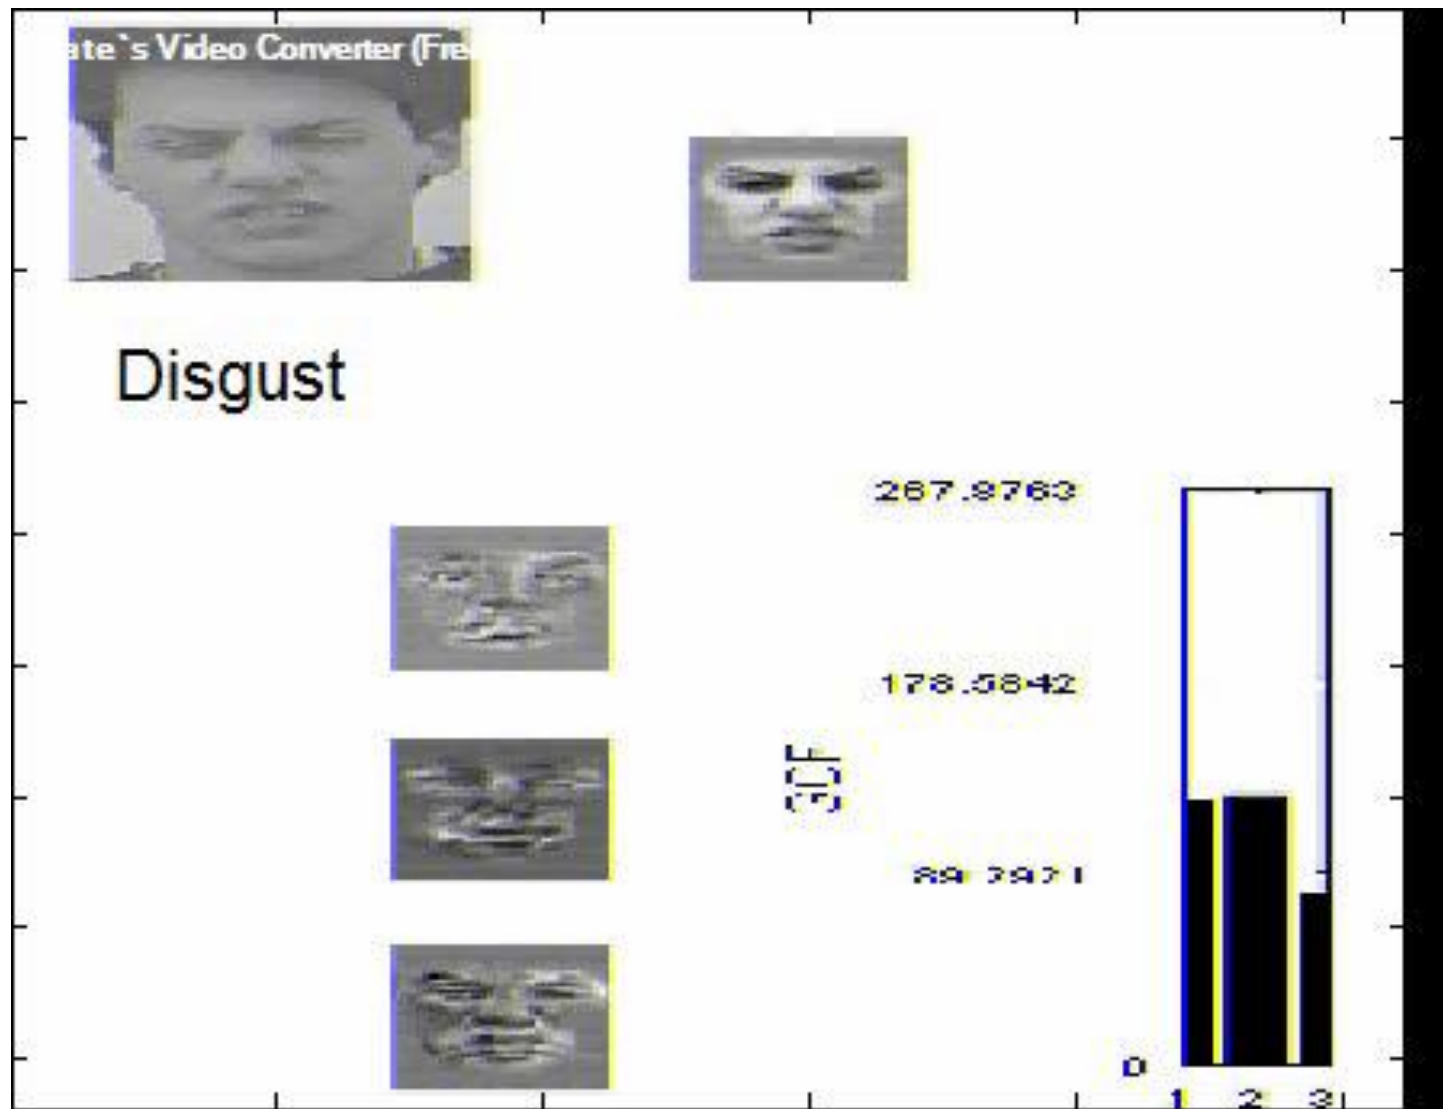# Happiness subspace

Disgust subspace

Facial Expression Recognition Movies

Facial Expression Recognition Movies

# Facial Expression Recognition Movies

Image Compression

Original Image

❑ Divide the original 372x492 image into patches:

   • Each patch is an instance that contains 12x12 pixels on a grid

❑ Consider each as a 144-D vector

L$_2$ error and PCA dim

# PCA compression: $144D \Rightarrow 60D$

PCA compression: 144D ⇒ 16D

PCA compression: 144D $\Rightarrow$ 6D

PCA compression: 144D ⇒ 3D

3 most important eigenvectors

60 most important eigenvectors

Looks like the discrete cosine bases of JPG!...

2D Discrete Cosine Basis

http://en.wikipedia.org/wiki/Discrete_cosine_transform

Noise Filtering

Noise Filtering

x   **x'**

U x

Noisy image

Denoised image
using 15 PCA components

PCA Shortcomings

PCA doesn't know about class labels!

PCA vs Fisher Linear Discriminant

- PCA maximizes variance, *independent of class*

  $\Rightarrow$ magenta

- FLD attempts to separate classes

  $\Rightarrow$ green line

Problematic Data Set for PCA

PCA cannot capture NON-LINEAR structure!

PCA Conclusions

❏ PCA
- finds orthonormal basis for data
- Sorts dimensions in order of "importance"
- Discard low significance dimensions

❏ Applications:
- Get compact description
- Remove noise
- Improve classification (hopefully)

❏ Not magic:
- Doesn't know class labels
- Can only capture linear variations

❏ One of many tricks to reduce dimensionality!

PCA Theory

Let $\mathbf{x} \in \mathbb{R}^N$

Let $\mathbf{X} = [\mathbf{x}_1, \ldots, \mathbf{x}_m] \in \mathbb{R}^{N \times m}$,
m: number of instances, N: dimension

Let $\mathbf{U} = \begin{pmatrix} \mathbf{u}_1^T \\ \vdots \\ \mathbf{u}_N^T \end{pmatrix} \in \mathbb{R}^{N \times N}$ orthogonal matrix, $\mathbf{U}\mathbf{U}^T = \mathbf{I}_N$

$$\mathbf{y} \doteq \mathbf{U}\mathbf{x}, \ \mathbf{x} = \mathbf{U}^T\mathbf{y} = \sum_{i=1}^N \mathbf{u}_i y_i$$

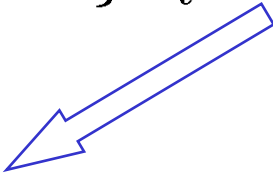$$\hat{\mathbf{x}} \doteq \sum_{i=1}^M \mathbf{u}_i y_i, \ (M \leq N)$$ approximation of $\mathbf{x}$ using M basis vectors only.

$$\varepsilon^2 \doteq \mathbb{E}\{\|\mathbf{x} - \hat{\mathbf{x}}\|^2\} = \frac{1}{m}\sum_{j=1}^m \|\mathbf{x}_j - \hat{\mathbf{x}}_j\|, \text{ average error}$$

GOAL:

$$\arg\min_{\mathbf{U}} \varepsilon^2, \text{ s.t } \mathbf{U}^T\mathbf{U} = \mathbf{I}_N$$

$$\varepsilon^2 \;=\; \mathbb{E}\{\|\mathbf{x} - \hat{\mathbf{x}}\|^2\} = \mathbb{E}\{\| \sum_{i=1}^{N} \mathbf{u}_i y_i - \sum_{i=1}^{M} \mathbf{u}_i y_i \|^2\}$$

$$=\; \mathbb{E}\{ \sum_{i=M+1}^{N} y_i \mathbf{u}_i^T \mathbf{u}_i y_i \} = \sum_{i=M+1}^{N} \mathbb{E}\{ y_i^2 \}$$

$$=\; \sum_{i=M+1}^{N} \mathbb{E}\{ (\mathbf{u}_i^T \mathbf{x})(\mathbf{x}^T \mathbf{u}_i) \}$$

$$=\; \sum_{i=M+1}^{N} \mathbf{u}_i^T \mathbb{E}\{\mathbf{x}\mathbf{x}^T\} \mathbf{u}_i \qquad \textbf{x} \text{ is centered!}$$

$$=\; \sum_{i=M+1}^{N} \mathbf{u}_i^T \mathbf{\Sigma} \mathbf{u}_i$$

Justification of Algorithm II

GOAL: $\arg\min\limits_{\mathbf{u}_{M+1},\dots,\mathbf{u}_N}\varepsilon^2$

Use Lagrange-multipliers for the constraints.

$$L = \varepsilon^2 - \sum_{i=M+1}^{N} \lambda_i(\mathbf{u}_i^T\mathbf{u}_i - 1)$$

$$= \sum_{i=M+1}^{N} \mathbf{u}_i^T\mathbf{\Sigma}\mathbf{u}_i - \sum_{i=M+1}^{N} \lambda_i(\mathbf{u}_i^T\mathbf{u}_i - 1)$$

$$\frac{\partial L}{\partial \mathbf{u}_i} = [2\mathbf{\Sigma}\mathbf{u}_i - 2\lambda_i\mathbf{u}_i] = 0$$

$$\frac{\partial L}{\partial \mathbf{u}_i} = [2\mathbf{\Sigma}\mathbf{u}_i - 2\lambda_i\mathbf{u}_i] = 0 \Rightarrow \mathbf{\Sigma}\mathbf{u}_i = \lambda_i\mathbf{u}_i$$

$$\Rightarrow [\mathbf{u}_i, \lambda_i] = \text{eigenvector/eigenvalue of } \mathbf{\Sigma}.$$

$$\varepsilon^2 = \sum_{i=M+1}^{N} \mathbf{u}_i^T \mathbf{\Sigma}\mathbf{u}_i = \sum_{i=M+1}^{N} \mathbf{u}_i^T \lambda_i \mathbf{u}_i = \sum_{i=M+1}^{N} \lambda_i$$

The error ε^2 is minimal

if $\lambda_{M+1}, \ldots \lambda_N$ are the smallest eigenvalues of $\mathbf{\Sigma}$,

and $\mathbf{u}_{M+1}, \ldots, \mathbf{u}_N$ are the corresponding eigenvectors.

Thanks for the Attention! ☺

Attic

Kernel PCA

Kernel PCA

Performing PCA in the feature space

Let $\mathbf{X} = [\mathbf{x}_1, \ldots, \mathbf{x}_m] \in \mathbb{R}^{N \times m}$,
m: number of instances, N: dimension

Lemma

$\mathbf{u}$ is eigenvector of $\Sigma \Rightarrow \mathbf{u}$ is a linear combinaton of the samples

Proof:

$$\lambda \mathbf{u} = \Sigma \mathbf{u} = \left(\frac{1}{m} \sum_{i=1}^{m} \mathbf{x}_i \mathbf{x}_i^T \right) \mathbf{u} = \frac{1}{m} \sum_{i=1}^{m} (\mathbf{x}_i^T \mathbf{u}) \mathbf{x}_i$$

$$\Rightarrow \mathbf{u} = \sum_{i=1}^{m} \underbrace{\frac{(\mathbf{x}_i^T \mathbf{u})}{\lambda m}}_{\alpha_i} \mathbf{x}_i = \sum_{i=1}^{m} \alpha_i \mathbf{x}_i$$

Kernel PCA

$$\mathbf{u} = \sum_{i=1}^{m} \underbrace{\frac{(\mathbf{x}_i^T \mathbf{u})}{\lambda m}}_{\alpha_i} \mathbf{x}_i = \sum_{i=1}^{m} \alpha_i \mathbf{x}_i \quad \mathbf{X} = [\mathbf{x}_1, \ldots, \mathbf{x}_m] \in \mathbb{R}^{N \times m},$$

Lemma

To calculate $\boldsymbol{\alpha} \in \mathbb{R}^m$

- just use inner products (Gram matrix): $K_{ij} = \mathbf{x}_i^T \mathbf{x}_j$
- don't need the actual values of $\mathbf{x}_i$

Kernel PCA

Proof

$$\Sigma \mathbf{u} = \lambda \mathbf{u}, \ \ \mathbf{u} = \sum_{j=1}^{m} \alpha_j \mathbf{x}_j$$

$$\Rightarrow \mathbf{x}_i^T \Sigma \mathbf{u} = \lambda \mathbf{x}_i^T \mathbf{u}$$

$$\Rightarrow \mathbf{x}_i^T \left(\frac{1}{m} \sum_{k=1}^{m} \mathbf{x}_k \mathbf{x}_k^T \right) \left(\sum_{j=1}^{m} \alpha_j \mathbf{x}_j \right) = \lambda \mathbf{x}_i^T \left(\sum_{j=1}^{m} \alpha_j \mathbf{x}_j \right)$$

$$\Rightarrow \frac{1}{m} \sum_{k=1}^{m} \sum_{j=1}^{m} (\mathbf{x}_i^T \mathbf{x}_k)(\mathbf{x}_k^T \mathbf{x}_j) \alpha_j = \lambda \sum_{j=1}^{m} (\mathbf{x}_i^T \mathbf{x}_j) \alpha_j$$

$$\Rightarrow \frac{1}{m} \mathbf{K}^2 \boldsymbol{\alpha} = \lambda \mathbf{K} \boldsymbol{\alpha} \quad \text{where } \mathbf{K} \in \mathbb{R}^{m \times m}$$

$$\Rightarrow \mathbf{K} \boldsymbol{\alpha} = m \lambda \boldsymbol{\alpha} \quad \text{If } \mathbf{K} \text{ is invertible (strictly pos def)}$$

Kernel PCA

❑ How to use α to calculate the projection of a new sample *t?*

$$\mathbf{u}^T\mathbf{t} = (\sum_{j=1}^{m} \alpha_j\mathbf{x}_j)^T\mathbf{t} = \sum_{j=1}^{m} \alpha_j K(\mathbf{x}_j, \mathbf{t})$$

Again, we don't need values of $\mathbf{x}_j$!

Let $K_{i,j} = \langle\phi(\mathbf{x}_i), \phi(\mathbf{x}_j)\rangle$

Where was I cheating? ☺

The data should be centered in the feature space, too!
But this is manageable...

$$\tilde{K}_{i,j} \doteq \left\langle \phi(\mathbf{x}_i) - \frac{1}{m}\sum_{k=1}^{m}\phi(\mathbf{x}_k), \quad \phi(\mathbf{x}_j) - \frac{1}{m}\sum_{k=1}^{m}\phi(\mathbf{x}_k) \right\rangle$$

Input points before kernel PCA

http://en.wikipedia.org/wiki/Kernel_principal_component_analysis

73

The three groups are distinguishable using the first component only $$k(\boldsymbol{x}, \boldsymbol{y}) = (\boldsymbol{x}^{\mathrm{T}}\boldsymbol{y} + 1)^2$$

Reconstructing… (Method B)

❑ … faster if train with…
- only people w/out glasses
- same lighting conditions